Absence of evidence is not evidence of absence
or lack of proof is not proof of lack
or how i couldn’t get it to work






google denied me

Retail

Gaming

Other:

What is your use case? *

Your answer

What kind of videos do you expect to send against the API?

- Social Media - Crowd-sourced content. Under 1 minute
- Social Media - Crowd-sourced content (> 1 minute)
- Professional Media - TV Shows, Movies
- News Content -
- Sports Content
- Surveillance Videos

What is the average length of the video content?

For Alpha - we will support Video content up to 1 hour.

- < 1 minute
- 1-10 minutes

scene dynamics hung … hard

Generating Videos with Scene Dynamics

This repository contains an implementation of Generating Videos with Scene Dynamics by Carl Vondrick, Hamed Pirsiavash, Antonio Torralba, to appear at NIPS 2016. The model learns to generate tiny videos using adversarial networks.

Example Generations

Below are some selected videos that are generated by our model. These videos are not real; they are hallucinated by a generative video model. While they are not photo-realistic, the motions are fairly reasonable for the scene category they are trained on.

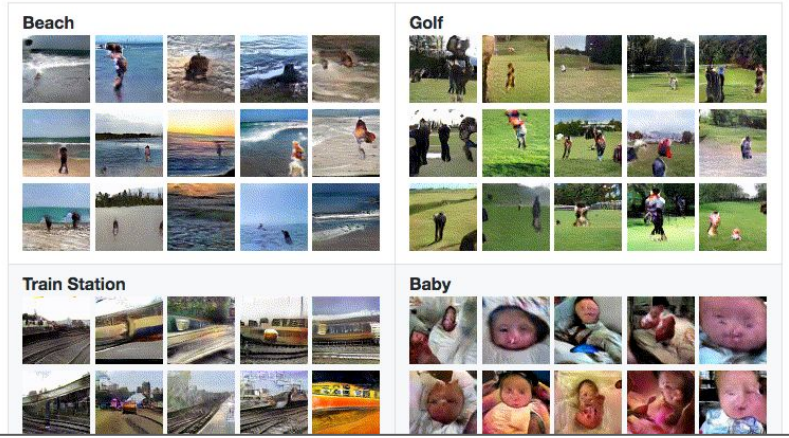

all the autoencoders i found were
really outdated

README.md

Learned Similarity Autoencoder for Modelling and Reconstructing Video Frames

Tesorflow implementation of Autoencoding beyond pixels using a learned similarity metric.

A lot of the architecture is derived from this codebase DCGAN-Tensorflow.

This project is designed to read 256x144 png's and that are indexed in numerial order.

This project was implemented for my project on reconstructing videos with neural networks - read more

To train a model with a dataset:

```
$ python main.py --dataset DATASETNAME --is_train True
```

You may want to adjust the amount of noise injected into the latent space:

```
$ python main.py --dataset DATASETNAME --is_train True --noise 0.5
```

This parameter controls the standard deviation of noise epsilon from mean of 0.

The output frames in sequence using an exisiting model:

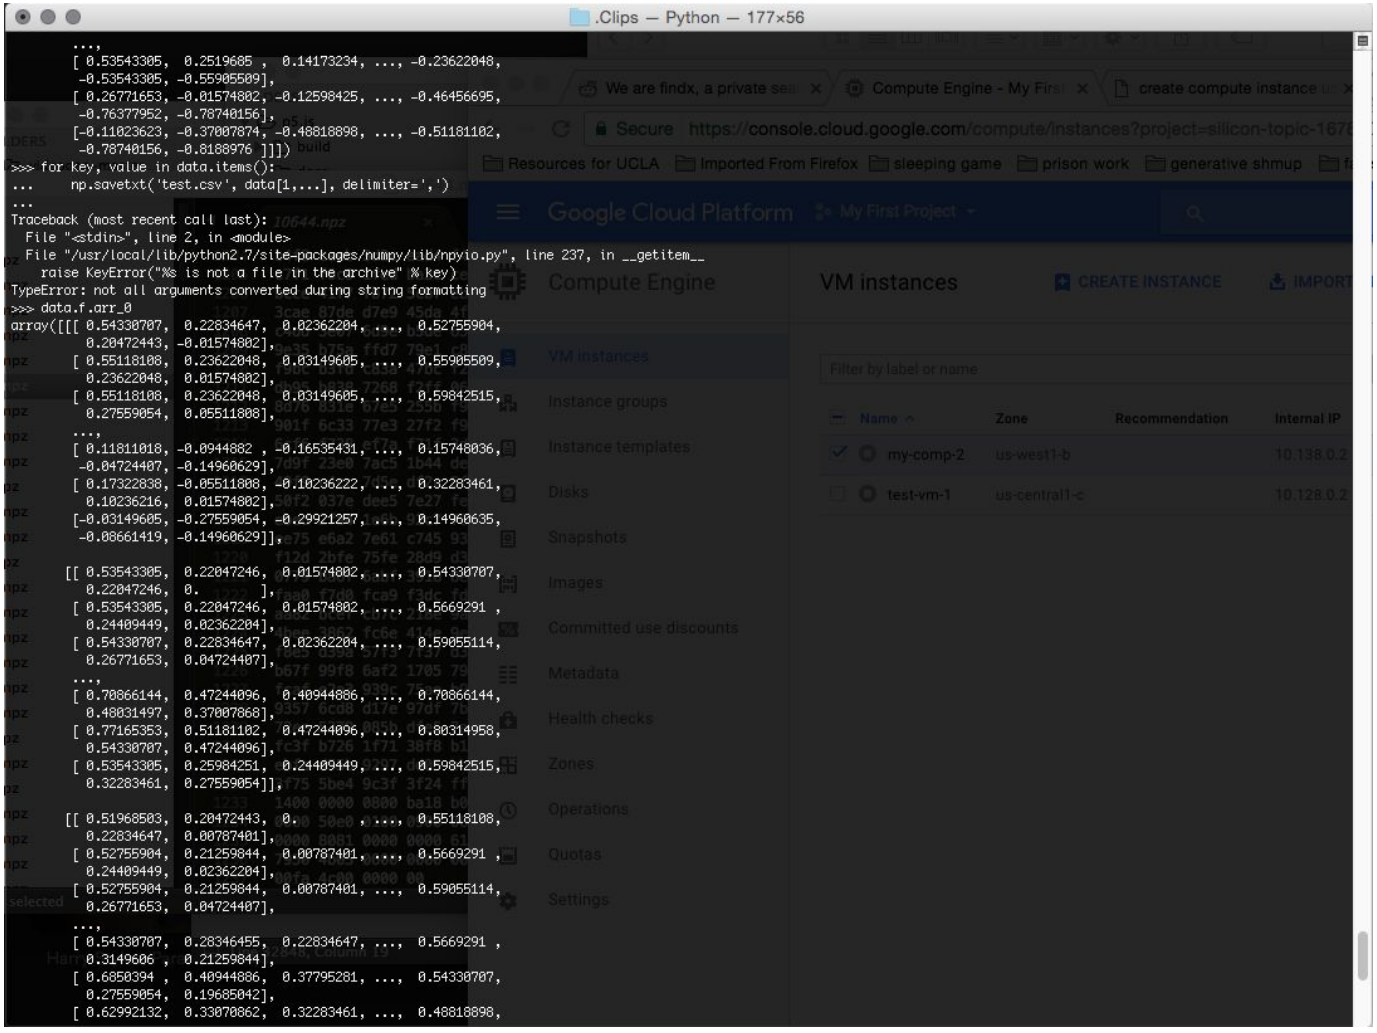

found some code to process the sequential images into 50k 32x32px 3D arrays in .npz format -- this is the only way to train on “full resolution” images

this one I got pretty far with but some of the code has to be rewritten for currently available version of tensorflow

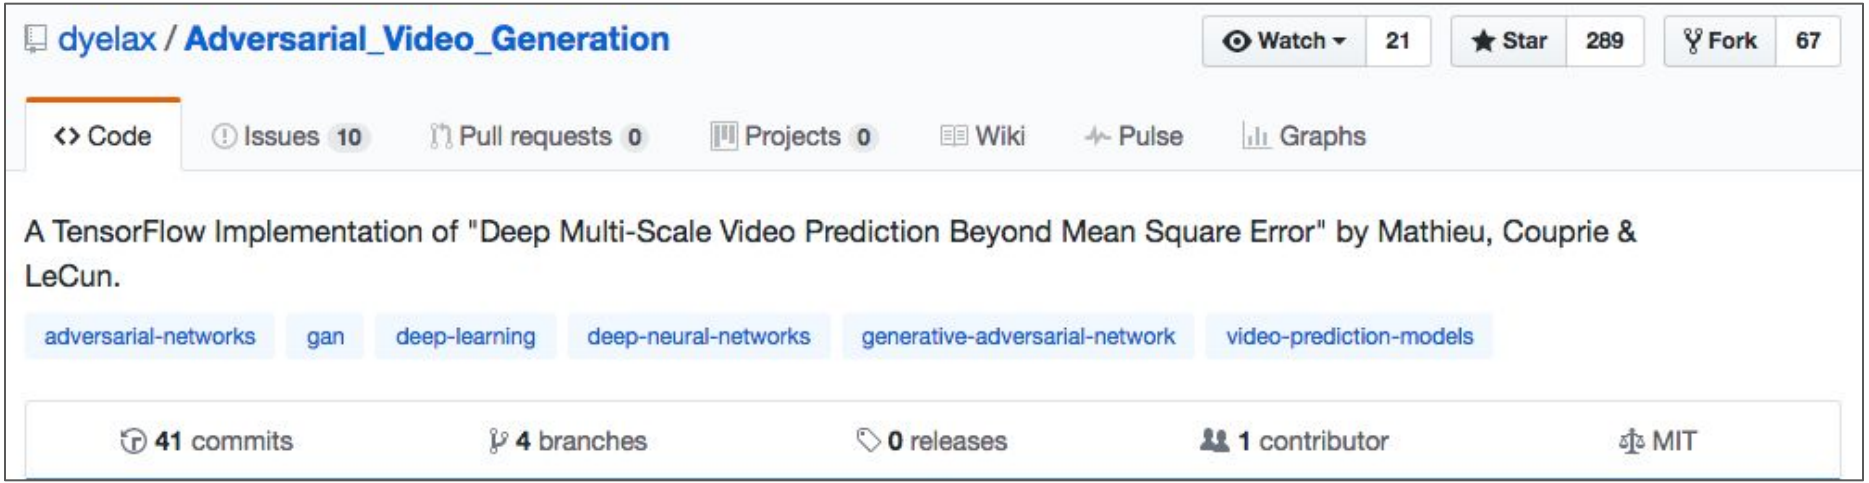

Next Steps?

try rewriting code for
dcgan or autoencoder?

or

find a secret awesome video generation slice of code that works effortlessly?